

# Topics to be covered

---

- Principles of I/O Hardware: I/O devices
- Device controllers
- Direct memory access
- Principles of I/O Software: Goals of Interrupt handlers
- Device drivers
- Device independent I/O software
- Secondary-Storage Structure: Disk structure
- Disk scheduling algorithm

# Components of I/O devices

- I/O devices have two components
  1. **Mechanical component**
  2. **Electronic component**
- The mechanical component is device itself.
- Electronic component of devices is called the **Device Controller**.



# Device Controller

---

- **Electronic component** which **controls the device**.
- It may handle multiple devices.
- There **may be more than one controller** per mechanical component (example: hard drive).
- Controller's tasks are:
  - It **converts serial bit stream to block of bytes**
  - **Perform error correction** if necessary
  - **Block of bytes is first assembled bit by bit** in buffer inside the controller
  - After verification, the block has been declared to be error free, and then it can be **copied to main memory**

# Memory-Mapped I/O

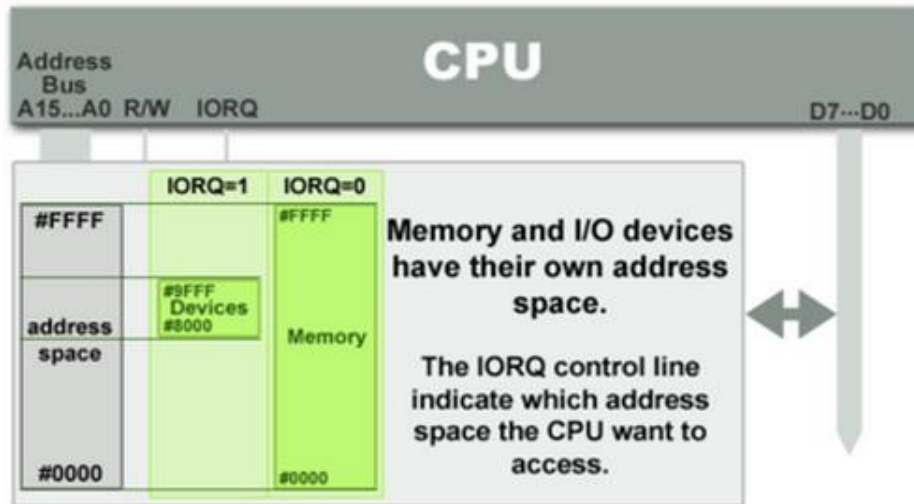
---

- Each device controller **has a few registers** that are **used for communicating with the CPU**.
- By **writing into these registers**, the OS can command the device to **deliver data, accept data, switch itself on or off**, or perform some action.
- By **reading from these registers** OS can learn **what the device's status is**, whether it is prepared to accept a new command and so on.
- There are two ways to communicate with control registers and the device buffers:
  1. I/O Port
  2. Memory mapped I/O

# I/O Port Vs Memory Mapped I/O


## I/O Port

- **Uses different address spaces** for memory and I/O devices
- **Uses a special class** of CPU instructions to access I/O devices



## Memory mapped I/O

- **Uses same address space** to address memory and I/O devices
- Access to the I/O devices using **regular instructions**



# Direct Memory Access

---

- Feature of computer systems that **allows certain hardware subsystems to access main memory (RAM), independent of the central processing unit (CPU).**
- Without DMA, when the CPU is using programmed input/output, it is typically fully occupied for the entire duration of the read or write operation, and is thus unavailable to perform other work.


# Direct Memory Access

---

- With DMA, the **CPU first initiates the transfer, then it does other operations while the transfer is in progress**, and it finally receives an interrupt from the DMA controller when the operation is done.
- This feature is **useful when the CPU needs to perform useful work while waiting** for a relatively slow I/O data transfer.
- Many hardware systems such as disk drive controllers, graphics cards, network cards and sound cards use DMA.

# Disk read-write without a DMA


- The disk controller **reads the block from the drive serially, bit by bit**, until the entire block is in the controller's buffer.
- Next, it **computes the checksum** to verify that no read errors have occurred.



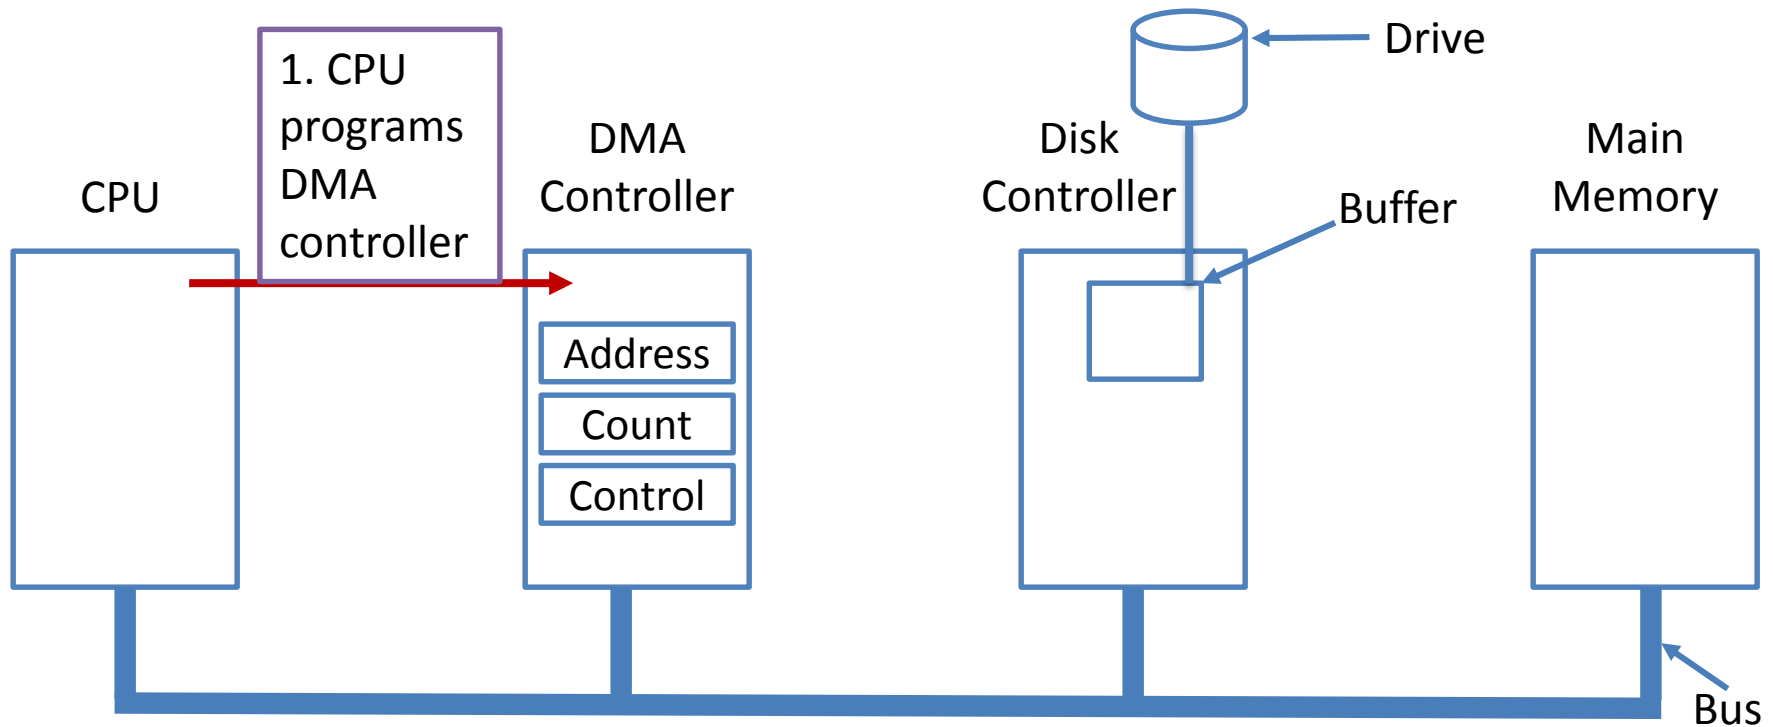


# Disk read-write without a DMA

- Then the **controller causes an interrupt**, so that OS can read the block from controller's buffer (a byte or a word at a time) by executing a loop.
- After reading every single part of the block from controller device register, the operating system will store them into the main memory.



# Disk read-write using DMA

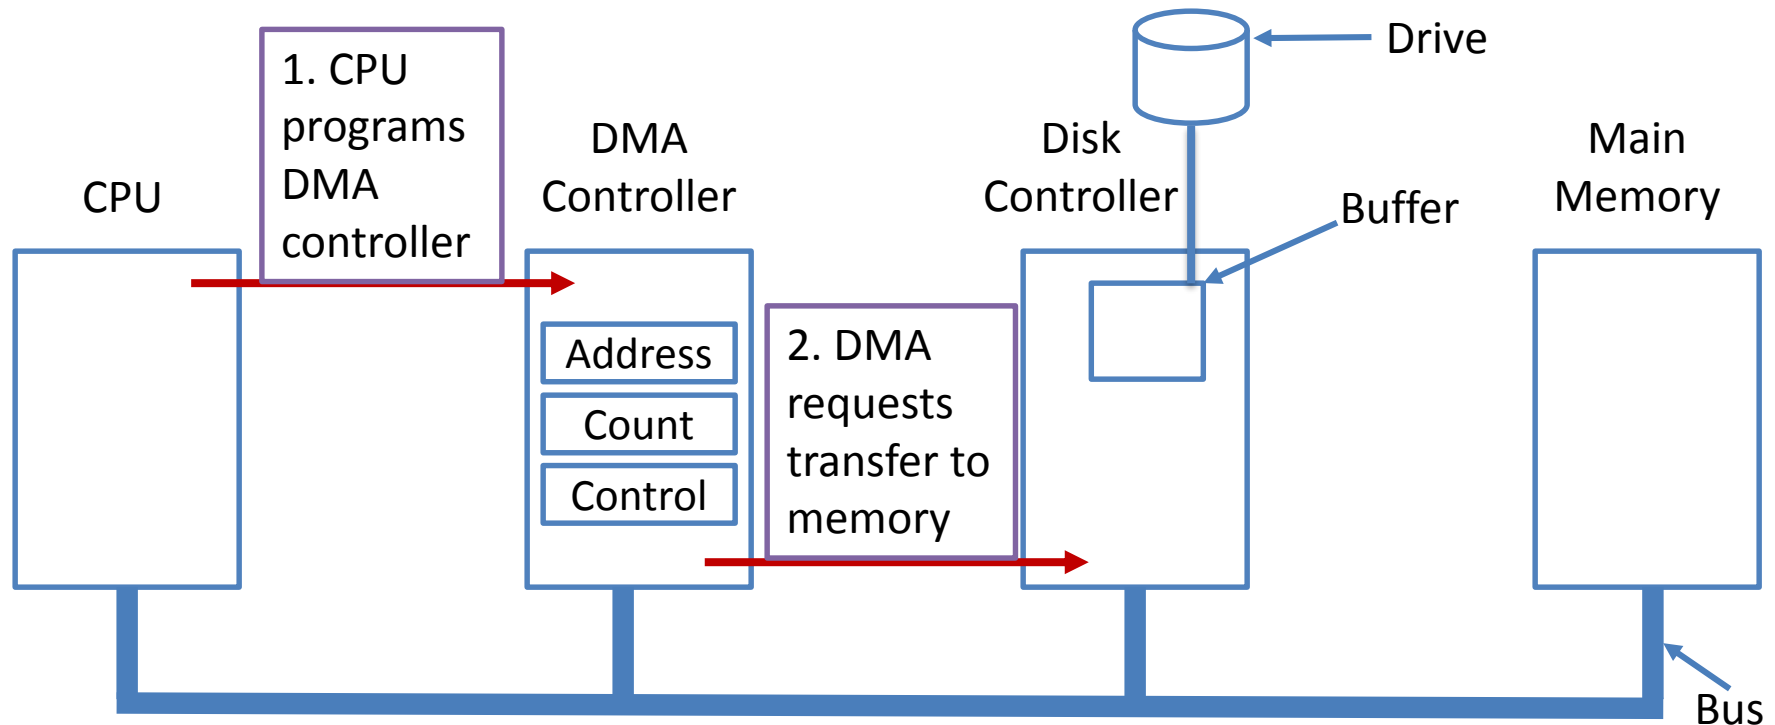


**Step 1:** First the **CPU programs the DMA controller** by setting its registers so it knows what to transfer where.

It also **issues a command to the disk controller** telling it to read data from the disk into its internal buffer and verify the checksum.

When valid data are in the disk controller's buffer, DMA can begin.

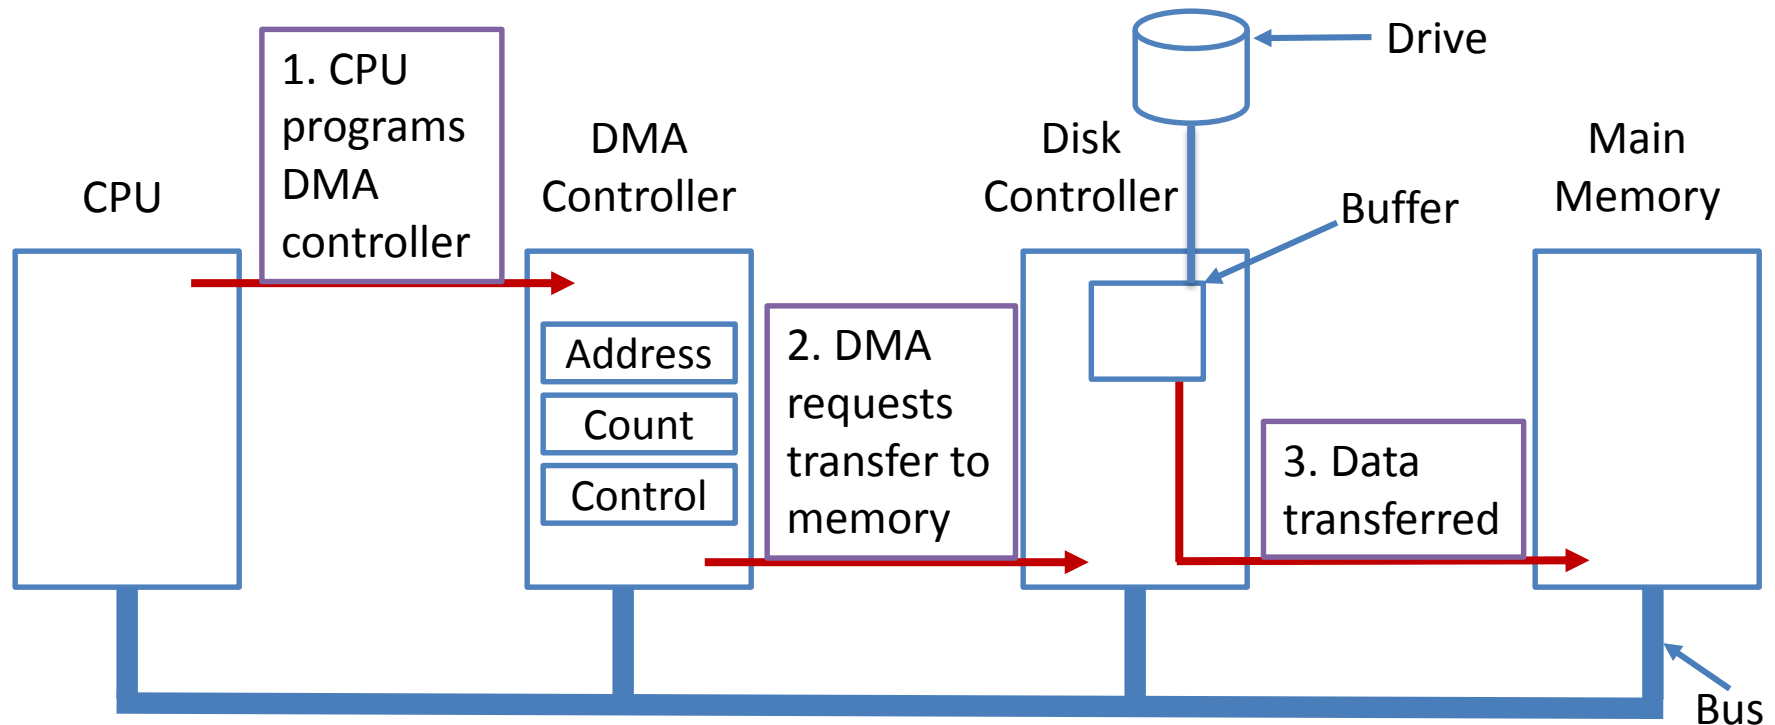
# Disk read-write using DMA



**Step 2:** The **DMA controller initiates the transfer** by issuing a read request over the bus to the disk controller.

This read request looks like any other read request, and the **disk controller does not know (or care)** whether it came from the CPU or from a DMA controller.

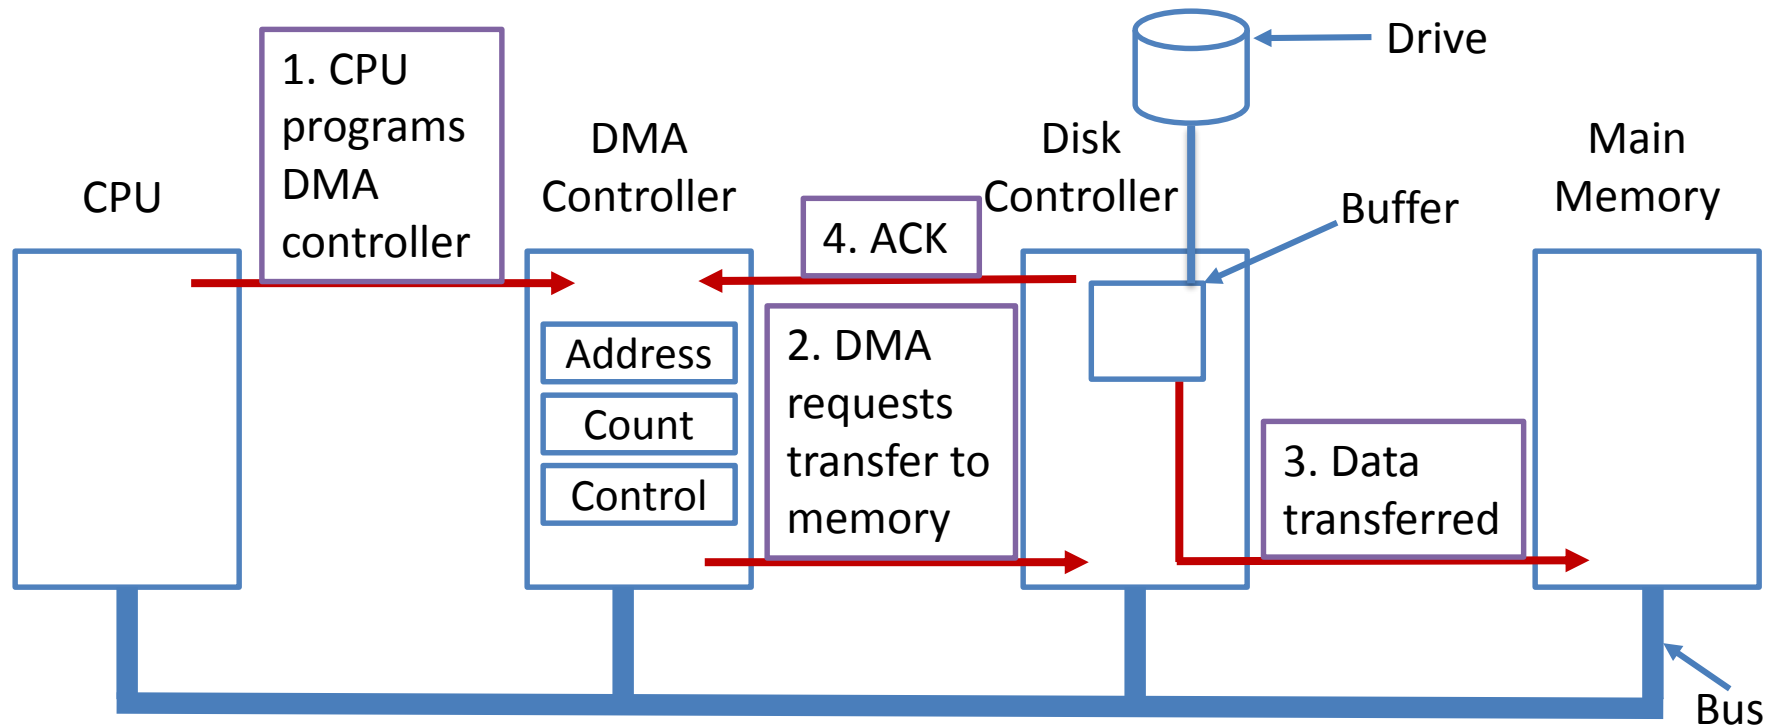
# Disk read-write using DMA



Typically, the memory address to write to is on the bus' address lines, so when the disk controller fetches the next word from its internal buffer, it knows where to write it.

**Step 3:** The write to memory is another standard bus cycle.

# Disk read-write using DMA




**Step 4:** When the write is complete, the **disk controller sends an acknowledgement** signal to the DMA controller, also over the bus.

The **DMA controller then increments the memory address** to use and decrements the byte count.

If the **byte count is still greater than 0**, steps 2 to 4 are repeated until it reaches 0.

# Disk read-write using DMA



At that time, the DMA controller interrupts the CPU to let it know that the transfer is now complete.

When the OS starts up, it does not have to copy the disk block to memory; it is already there.

# Goals of I/O Software

---

## 1. Device independence

- It should be possible to write **programs that can access any I/O devices** without having to specify device in advance.
- For example, a program that reads a file as input should be able to read a file on a floppy disk, on a hard disk, or on a CD-ROM, without having to modify the program for each different device.

## 2. Uniform naming

- Name of file or device should be some specific string or number. It **must not depend upon device** in any way.
- All files and devices are addressed the same way: by a path name.

# Goals of I/O Software

---

## 3. Error handling

- **Error should be handled as close to hardware** as possible.
- If any controller generates error then it tries to solve that error itself. If controller can't solve that error then device driver should handle that error, perhaps by reading all blocks again.
- Many times when error occurs, error is solved in lower layer. If lower layer is not able to handle error then problem should be told to upper layer.
- In many cases error recovery can be done at a lower layer without the upper layers even knowing about error.



# Goals of I/O Software

---

## 4. Synchronous vs. asynchronous transfers

- Most of devices are asynchronous device. CPU starts transfer and goes off to do something else until interrupt occurs.
- I/O Software **needs to support both the types of devices.**

## 5. Buffering

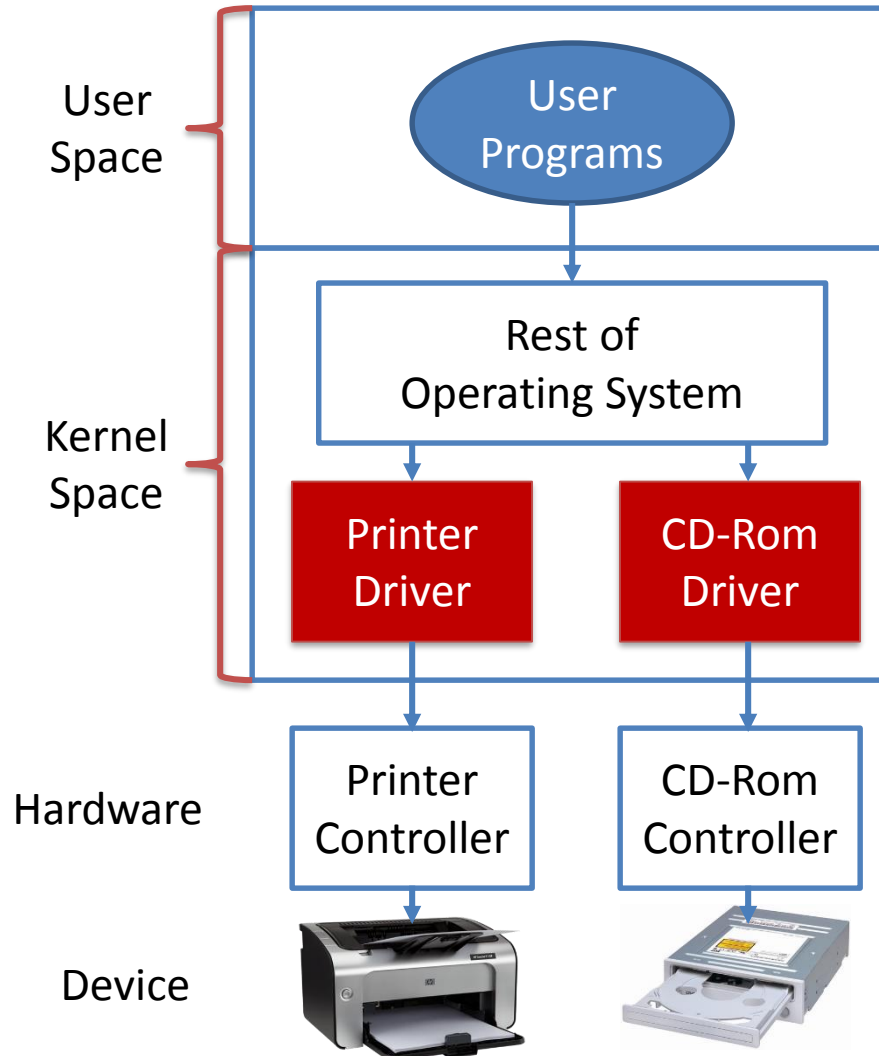
- **Data comes in main memory cannot be stored directly.**
- For example data packets come from the network cannot be directly stored in physical memory.
- Packets have to be put into output buffer for examining them.
- Some devices have several real-time constraints, so data must be put into output buffer in advance to decouple the rate at which buffer is filled and the rate at which it is emptied, in order to avoid buffer under runs.

# Device driver

---

- I/O devices which are plugged with computer have **some specific code for controlling them**. This code is called the **device driver**.
- Each device driver normally handles one device type, or at most one class of closely related devices.
- Generally device driver is delivered along with the device by device manufacturer.
- Device drivers are normally positioned below the rest of Operating System.

# Logical positioning of device drivers



# Functions of device drivers

---

1. Device driver **accept abstract read and write requests** from device independent software.
2. Device driver must **initialize the device** if needed.
3. It also **controls power requirement** and **log event**.
4. It also **checks statues of devices**. If it is currently in use then queue the request for latter processing. If device is in idle state then request can be handled now.
5. Pluggable device can be added or removed while the computer is running. At that time the device driver **inform CPU** that the user has suddenly removed the device from system.


# RAID

---

- RAID (**Redundant Array of Independent Disks**)
- RAID is a **data storage virtualization technology** that **combines multiple physical disk drive components** into a single logical unit for the **purposes of data redundancy, performance improvement, large storage capacity** or all.
- Data is distributed across the drives in one of several ways, referred to as RAID levels, depending on the required level of redundancy and performance.
- All RAID have the property that the **data are distributed over drives, to allow parallel operation.**
- There are **7 levels of RAID.**

# RAID 0

- It **splits data (file) into blocks of data**.
- **Stripe the blocks** across disks in the system.
- In the diagram to the right, the **odd blocks are written to disk 1** and the **even blocks to disk 2**.




# RAID 0

---

- It is **easy to implement**.
- **No parity calculation** overhead is involved.
- It **provides good performance** by spreading the load across many channels and drives.
- It **provides no redundancy** or **error detection**.
- Not true RAID because there is no fault tolerance. The **failure of just one drive will result in all data in an array being lost**.
- After certain amount of drives, **performance does not increase significantly**.
- It **requires minimum 2 drives** to implement.

# RAID 1

- A **complete file is stored on a single disk.**
- A **second disk contains an exact copy of the file.**
- It **provides complete redundancy of data.**
- Read performance can be improved
  - same file data can be read in parallel
- Write performance suffers
  - must write the data out twice





# RAID 1

---

- Most **expensive RAID implementation**.
- It **requires twice as much storage space**.
- In case a **drive fails, data do not have to be rebuild**, they just have to be copied to the replacement drive.
- The main disadvantage is that the **effective storage capacity is only half of the total drive capacity** because all data get written twice.


# RAID 1

---

- Software RAID 1 solutions do not always allow a hot swap of a failed drive.
- That means the failed drive can only be replaced after powering down the computer it is attached to.
- For servers that are used simultaneously by many people, this may not be acceptable. Such systems typically use hardware controllers that do support hot swapping.

# RAID 2

- It **stripes data at bit level (rather than block level) with dedicated Hamming-code parity**.
- It **uses ECC (Error Correcting Code)** to monitor correctness of information on disk.
- A **parity disk is then used to reconstruct corrupted or lost data**.




# RAID 2

---

- Imagine **splitting each byte into a pair of 4-bit nibbles**, then **adding Hamming code to each one to form a 7-bit word**, of which **bit 1, 2 and 4** were **parity bits**.
- In this RAID level 2 **each of seven drives needs synchronized** in terms of arm position and rotational position, and then it would be possible to write the 7-bit Hamming coded word over the seven drives.
- Here, **losing one drive did not cause problem**, which can be handled by Hamming code on the fly.
- Big problem is performance
  - must have to **read data plus ECC code** from other disks
  - for a **write, must have to modify data, ECC, and parity disks**


# RAID 3

- This technique **uses striping** and **dedicates one drive for storing parity** information.
- Here **single parity bit is computed for each data word** and written to a parity drive.
- The embedded ECC information is used to detect errors.
- As in RAID level 2 the drives must be exactly synchronized.



# RAID 4

- This level **uses large stripes (block level striping)**, which means you can read records from any single drive.
- They **do not require synchronization of drives.**



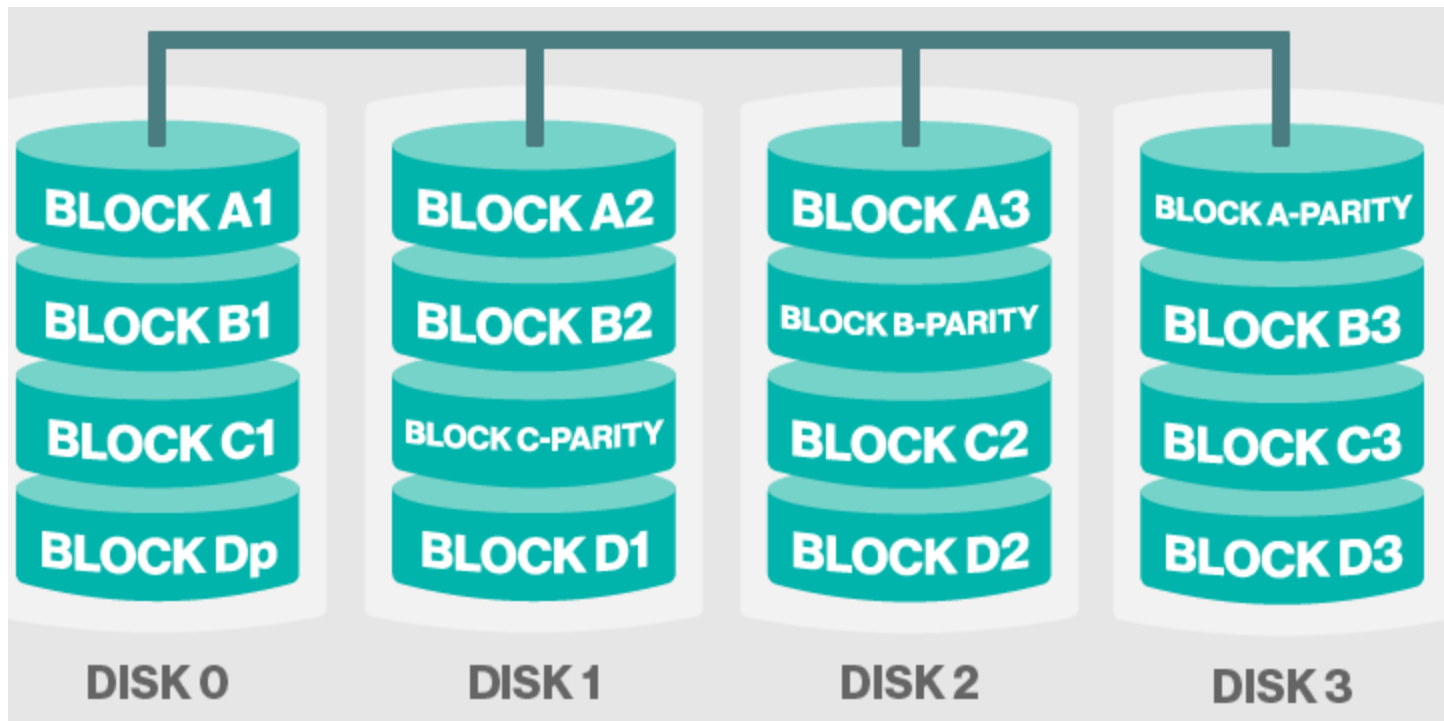
# RAID 4

---

- RAID level 4 is like RAID level 0, with **strip-for-strip parity written onto an extra drive**, for example, if each strip is k bytes long, all strips are EXCLUSIVE ORed together, resulting in a parity strip k bytes long.
- If a **drive crashes, the lost bytes can be recomputed from the parity drive** by reading the entire set of drives.
- This design protects against the loss of a drive but performs poorly for small updates, if one sector is changed, it is necessary to read all the drives in order to recalculate the parity.
- It **creates heavy load** on parity drive.

# RAID 5

- This level is **based on block-level striping with parity**.
- The **parity information is striped across each drive**.






# RAID 5

---


- As with **RAID level 4**, there is a **heavy load in the parity drive**, it **may become bottleneck**.
- This **bottleneck can be eliminated in RAID level 5 by distributing the parity bits uniformly** over all the drives.

## Architecture-

- The entire disk is divided into **platters**.
- Each platter consists of concentric circles called as **tracks**.
- These tracks are further divided into **sectors** which are the smallest divisions in the disk.




**Disk divided into tracks**



**Track divided into sectors**


- A **cylinder** is formed by combining the tracks at a given radius of a disk pack.



- There exists a mechanical arm called as **Read / Write head**.
- It is used to read from and write to the disk.

# Definitions

- Seek time: The time to move the arm to the proper cylinder.
- Rotational delay: The time for the proper sector to rotate under the head.



# Disk Arm Scheduling Algorithm

---

- Various types of disk arm scheduling algorithms are available to decrease mean seek time.
  1. FCFS (First come first serve)
  2. SSTF (Shorted seek time first)
  3. SCAN
  4. C-SCAN
  5. LOOK (Elevator)
  6. C-LOOK


# Example for Disk Arm Scheduling Algorithm

---

- Consider an imaginary disk with 51 cylinders. A request comes in to read a block on cylinder 11. While the seek to cylinder 11 is in progress, new requests come in for cylinders 1, 36, 16, 34, 9, and 12, in that order.
- Starting from the current head position, what is the total distance (in cylinders) that the disk arm moves to satisfy all the pending requests, for each of the following disk scheduling Algorithms?

# FCSC (First come first serve)

- Here **requests are served in the order of their arrival.**




1,  
36,  
16,  
34,  
9,  
12

- Disk movement will be 11, 1, 36, 16, 34, 9 and 12.
- Total cylinder movement:  $(11-1) + (36-1) + (36-16) + (34-16) + (34-9) + (12-9) = 111$

# SSTF (Shortest seek time first)

- We can minimize the disk movement by **serving the request closest to the current position** of the head.



- Disk movement will be 11, 12, 9, 16, 1, 34, 36.
- Total cylinder movement:  $(12-11) + (12-9) + (16-9) + (16-1) + (34-1) + (36-34) = 61$




# LOOK (Elevator)

---

- **Keep moving in the same direction until there are no more outstanding requests pending in that direction, then algorithm switches the direction.**
- After switching the direction the arm will move to handle any request on the way. Here **first go it moves in up direction then goes in down direction.**
- This is also called as **elevator algorithm.**
- In the elevator algorithm, the **software maintains 1 bit: the current direction bit, which takes the value either UP or DOWN.**

# LOOK (Elevator)



1,  
36,  
16,  
34,  
9,  
12


- Disk movement will be 11, 12, 16, 34, 36, 9, 1.
- Total cylinder movement:  $(12-11) + (16-12) + (34-16) + (36-34) + (36-9) + (9-1)=60$

# C-LOOK

---

- Keep moving in the same direction until there are no more outstanding requests pending in that direction, then algorithm switches direction.
- When switching occurs the arm goes to the lowest numbered cylinder with pending requests and from there it continues moving in upward direction again.

# C-LOOK



1,  
36,  
16,  
34,  
9,  
12


- Disk movement will be 11, 12, 16, 34, 36, 1, 9.
- Total cylinder movement:  $(12-11) + (16-12) + (34-16) + (36-34) + (36-1) + (9-1) = 68$

# SCAN

---

- From the current position disk arm starts in up direction and moves towards the end, serving all the pending requests until end.
- At that end arm direction is reversed (down) and moves towards the other end serving the pending requests on the way.

# SCAN



1,  
36,  
16,  
34,  
9,  
12


- Disk movement will be 11, 12, 16, 34, 36, 50, 9, 1.
- Total cylinder movement:  $(12-11) + (16-12) + (34-16) + (36-34) + (50-36) + (50-9) + (9-1) = 88$

# CSCAN

---

- From the **current position disk arm starts in up direction and moves towards the end, serving request until end.**
- At the end the arm **direction is reversed (down), and arm directly goes to other end and again continues moving in upward direction.**

# CSCAN



- Disk movement will be 11, 12, 16, 34, 36, 50, 0, 1, 9.
- Total cylinder movement:  $(12-11) + (16-12) + (34-16) + (36-34) + (50-36) + (50-0) + (1-0) + (9-1) = 98$



# Examples for Disk Arm Scheduling Algorithm

---

- Consider an imaginary disk with 45 cylinders. A request comes in to read a block on cylinder 20. While the seek to cylinder 20 is in progress, new requests come in for cylinders 10, 22, 20, 2, 40, 6, and 38 in that order.
- Starting from the current head position, what is the total distance (in cylinders) that the disk arm moves to satisfy all the pending requests and how much seek time is needed for, for each of the disk scheduling algorithms if a seek takes 6 msec per cylinder moved?