

Aayush Shah

19BCE245

1 May 2021

Practical 9

Linear regression

- **Definition :** Write a program to implement simple and multiple linear regression.

Use the model parameters to make prediction for new data. Evaluate the model on standard evaluation measures. Your program should compute and display important model statistics (all the ones which are reported by statsmodels) and critically comment on these statistics.

- **Code :**

- **Simple Linear Regression**

```
#!/usr/bin/env python3
```

```
import random
```

```
import matplotlib.pyplot as plt
```

```
#user input
```

```
size_of_data = int(input("Enter size of data you wanted : "))
```

```
starting_range = int(input("Enter starting limit of data : "))
```

```
ending_range = int(input("Enter ending limit of data : "))
```

```
#random data generating
```

```
independent_data =
```

```
random.sample(range(starting_range,ending_range),size_of_data)
```

```
dependent_data =
```

```
random.sample(range(starting_range,ending_range),size_of_data)
```

```
#sorting data
```

```
independent_data.sort()
```

```
dependent_data.sort()
```

```
x = independent_data
```

```
y = dependent_data
```

```

#calculating intercept and slope
n=len(x)
E_y=sum(y)
E_x=sum(x)
E_xy=sum([i*j for i,j in zip(x,y)])
E_x2=sum([i*i for i in x])
E_y2=sum([i*i for i in y])
intercept=(E_y*E_x2-E_x*E_xy)/(E_x2-E_x**2)
slope=(n*E_xy-E_x*E_y)/(n*E_x2-E_x**2)

#generating graph
x_val=list(range(80));
y_val=[intercept+slope*i for i in x_val];
print(f"\nIntercept = {round(intercept,2)}\nSlope =
{round(slope,2)}")
plt.plot(x,y)
plt.plot(x_val,y_val)
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.title("Linear Regression")
plt.show()

```

• Sample I/O :

```

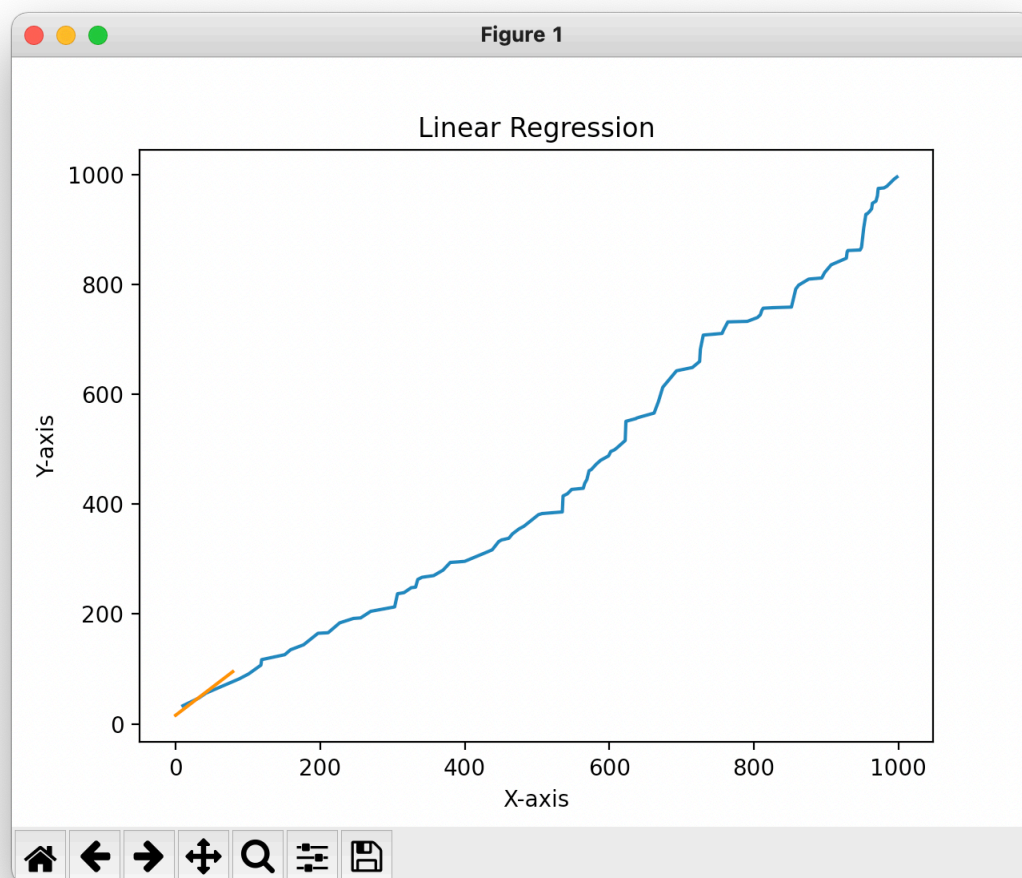
Enter size of data you wanted : 100
Enter starting limit of data : 1
Enter ending limit of data : 1000

Intercept = 15.71
Slope = 1.0

```

Running... CPU 0% Memory 96.4M Symbol ↕ Tabs: 4 ↕ Line 25, Column 7

• **Generated Graph :**



• **Multiple Linear Regression**

```

1. #on boston house-prices dataset
2. #Samples Total 506
3. #Dimensionality 13
4. #Features real, positive
5. #targets real
6. import numpy as np
7. from sklearn import datasets, linear_model, metrics
8. from sklearn.preprocessing import StandardScaler
9. # Load the boston dataset
10.X, y = datasets.load_boston(return_X_y=True)
11.X_train=X[0:400,:]
12.y_train=y[0:400]
13.X_test=X[400:506,:]
14.y_test=y[400:506]
15.

```

```
16.print(X_train)
17.print(X_test)
18.
19.scaler = StandardScaler()
20.scaler.fit(X_train)
21.X_train = scaler.transform(X_train)
22.X_test=scaler.transform(X_test)
23.
24.# Create linear regression object
25.regr =
    linear_model.SGDRegressor(max_iter=1000,learning_rate='constant', eta0=0.01)
26.
27.# Train the model using the training sets
28.regr.fit(X_train, y_train)
29.predictions=regr.predict(X_test)
30.
31.print("Predictions :",predictions)
32.
33.# Calculating the mean squared error of the prediction
34.mse = np.sum((predictions-y_test)**2)/len(y_test)
35.print(f'Mean Square Error : {mse}')
36.
37.# Calculating the mean absolute error of the prediction
38.mae = np.sum(abs(predictions-y_test))/len(y_test)
39.print(f'Mean Absolute Error : {mae}')
40.
41.print("Predicted Coefficients : ")
42.
43.# Calculating the slope
44.for i in range(len(regr.coef_)):
45.    print(f'x{i+1} : {(regr.coef_)[i]}')
46.
47.print(f'Predicted Intercept : {(regr.intercept_)[0]}')
```

• Sample I/O :

```
[6.32000e-03 1.80000e+01 2.31000e+00 ... 1.53000e+01 3.96900e+02
4.98000e+00]
[2.73100e-02 0.00000e+00 7.07000e+00 ... 1.78000e+01 3.96900e+02
9.14000e+00]
[2.72900e-02 0.00000e+00 7.07000e+00 ... 1.78000e+01 3.92830e+02
4.03000e+00]
...
```

...

```
Predictions : [15.39973066 21.91917666 22.40359965 16.3305633
10.17980293 11.46469693
10.17967746 23.65286405 16.54062672 23.62127872 18.14757829
20.25744476
```

...

```
27.08610038
26.56773696 30.11068853 19.65948144 19.53006282 24.68588755
14.42850297
23.1258666 25.88602494 27.4429049 31.67099326 33.44514723
24.57899122
22.95310467 26.36491009 23.69053448 24.97576848 12.09668802
8.12024726
3.02824115 14.454009 16.72423774 20.17274638 20.11266642
15.93373159
13.21669587 18.9172974 21.16565039 17.98863125 20.46322271
23.81461123
22.40055793 28.57174142 26.91965471 22.44940593]
```

Mean Square Error : 36.2017819105814

Mean Absolute Error : 5.155303536458874

Predicted Coefficients :

x1 : -0.6683920245155992

x2 : 1.099755826411337

x3 : 0.5819253390429671

x4 : 0.589416098097974

x5 : -1.6779094099166818

x6 : 3.323290268510643

x7 : 0.35664291244947427

x8 : -2.77481020793202

x9 : 3.0825463439951513

x10 : -1.7275294645889065

x11 : -1.538601619181304

x12 : 0.4569558461354575

x13 : -3.855921650421727

Predicted Intercept : 24.173442899588068

Run Succeeded Time 1 662 ms Peak Memory 67.1M Symbol Spaces: 4 Line 47, Column 55